



# Filesystems — Interface


---

MARCH 27, 2025

Sara McAllister  
[sjmcalli@cs.cmu.edu](mailto:sjmcalli@cs.cmu.edu)  
[saramcallister.github.io](https://saramcallister.github.io)

# Situating this lecture – Operating Systems (ECE 344)

Date and recording	Lecture Note
1/9	Introduction
1/10	Architecture Support
1/16	Architecture Support
1/17	Architecture Support
1/23	Processes
1/24	Processes
1/30	Threads
1/31	Threads
2/6	Threads, Synchronization (I)
2/7	Synchronization (I), Synchronization (II)
2/13	Synchronization (II)
2/14	Synchronization (II)
2/28	Memory Management (I)
3/6	Memory Management (I)
3/7	Memory Management (I)
3/14	Memory Management (II): Paging
3/20	Memory Management (II): Paging
3/21	Memory Management (III): Replacement



# What we've covered so far:

---

Processes

Threads


Storage Devices



Hard disk drives

Cycle-head-sector (CHS) addressing

Logical block addresses (LBAs)




0 1 2 3 4



How can we abstract storage devices?

# How do we abstract persistent storage?

Filesystems: a system for organizing, managing, and accessing *files* and *directories* stored on persistent media



Local filesystems (UNIX)



File sharing platforms





# Why do we need filesystems?

---

- Durability across restarts and crashes
- Naming and organization
- Sharing data between processes and users



# What makes filesystems challenging?

---

- Crash recovery (next lecture)
- **Performance + concurrency**
- **Sharing + security**



# What we're covering today: Filesystem interface

---

What are files?

`open` and `fds`

Directories and ownership



# What are files?

---

A file is a sequence of bytes, logically grouped together

- Allow users/ processes to access all files the same way

In UNIX, “Everything is a file”

- `/dev/sda2` (disk partition)
- `/dev/tty2` (terminal)
- `/proc/$$/status` (current process status)





# Typical file metadata

---

**Name** - 14 characters? 8.3? 255?

**Identifier** - “file number” (usually internal)

Size - two meanings (next lecture)

**Protection** - Who can do what?

Time, date, last modifier - monitoring, cleaning up



# Operations on Files

---

**Write, Read - often via position pointer/ "cursor"**

**Seek - adjust position pointer for next access**

Append - write at end of file (implicit synchronization)

Rename - change name of file inside directory and  
(maybe) move a file between 2 directories



# I/O to a file

---

Need to read and write to a file

So how do you read and write to one?

Try 1: `read("example.txt", buffer, num_bytes);`

- What if the file is large?

# I/O to a file

---

Need to read and write to a file

So how do you read and write to one?

Try 1: `read("example.txt", buffer, num_bytes);`

- What if the file is large?

Try 2: `read("example.txt", buffer, num_bytes, start_loc);`

- What is notably inefficient about this?
- What's the solution?





# What we're covering today: Filesystem interface

---

What are files?

`open` and `fds`

Directories and ownership



## Solution: Open-file state

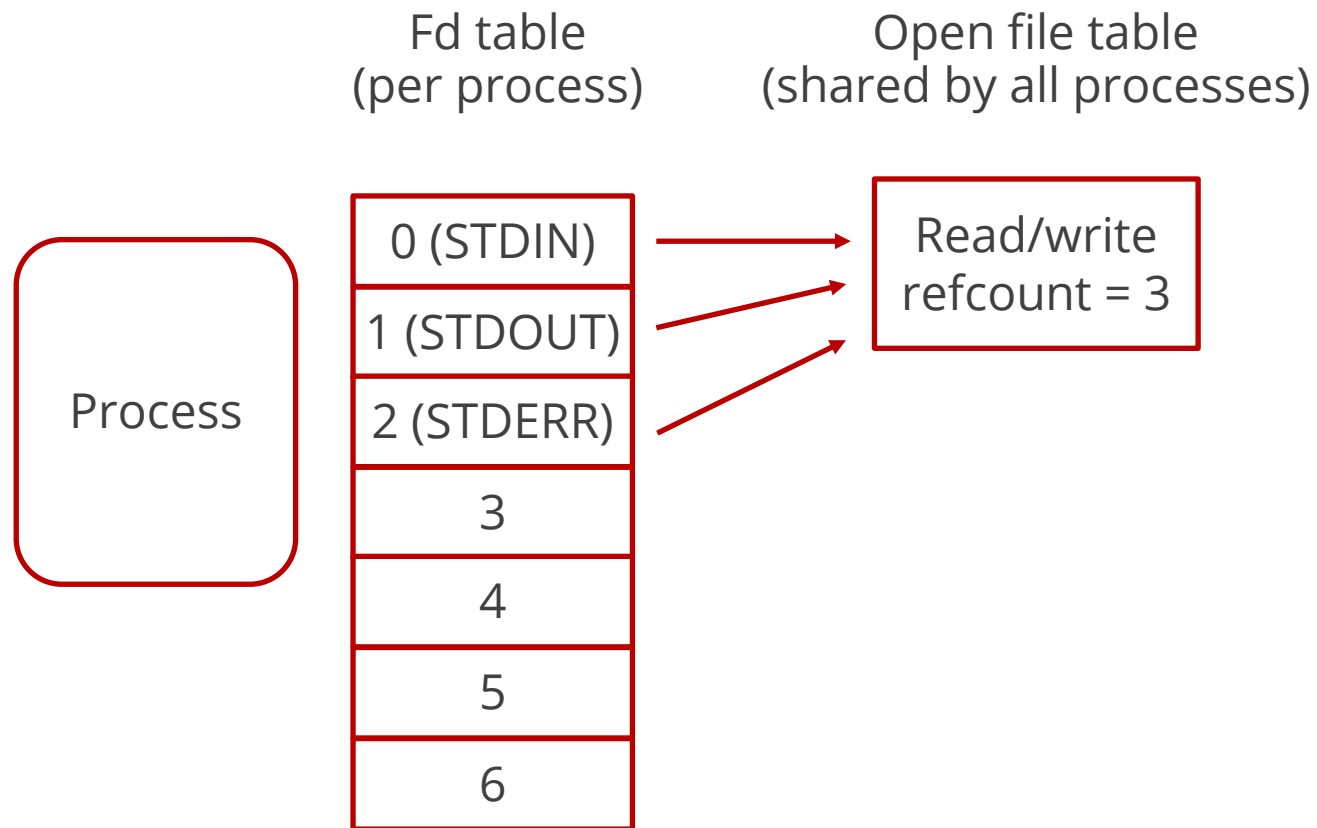
---

Add an `open` operation + “state”

“Open-file” structure (file descriptor) stores

- Filesystem / partition
- Filesystem-relative file number
- Operations allowed: eg read vs write
- Cursor position

# Open files (UNIX model)





# Open files (Unix model)

---

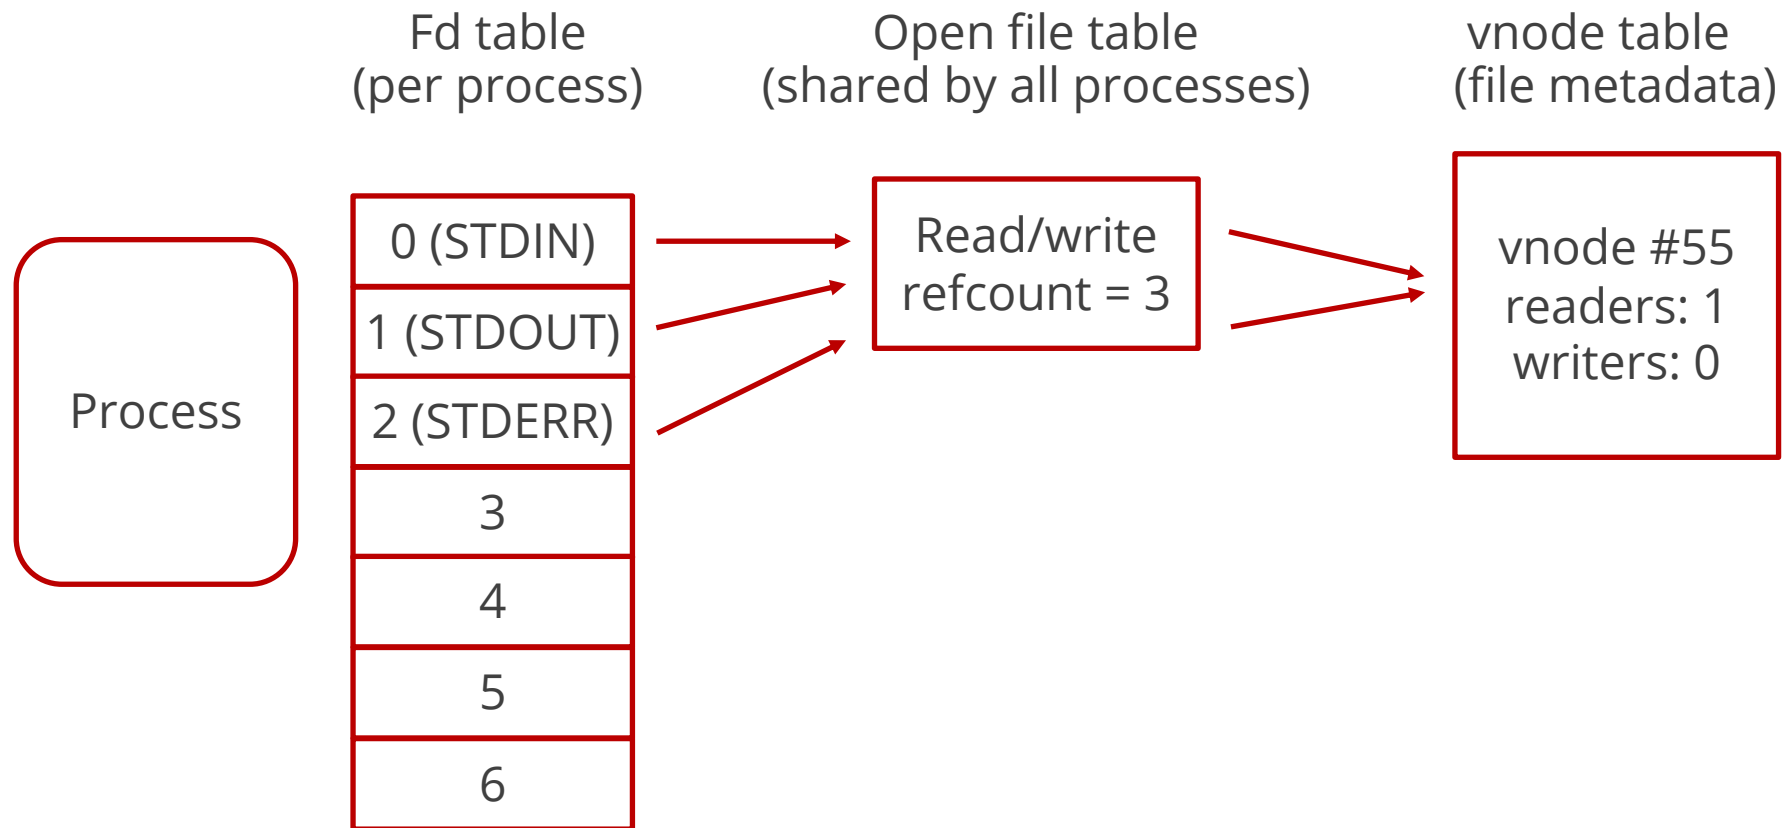
“In-core” file state - avoid going to disk repeatedly

- Mirror of on-disk structure (File number, size, permissions, modification time, ...)
- Any other filesystem specific information

Shared when file is opened multiple times



# Open files (UNIX model)



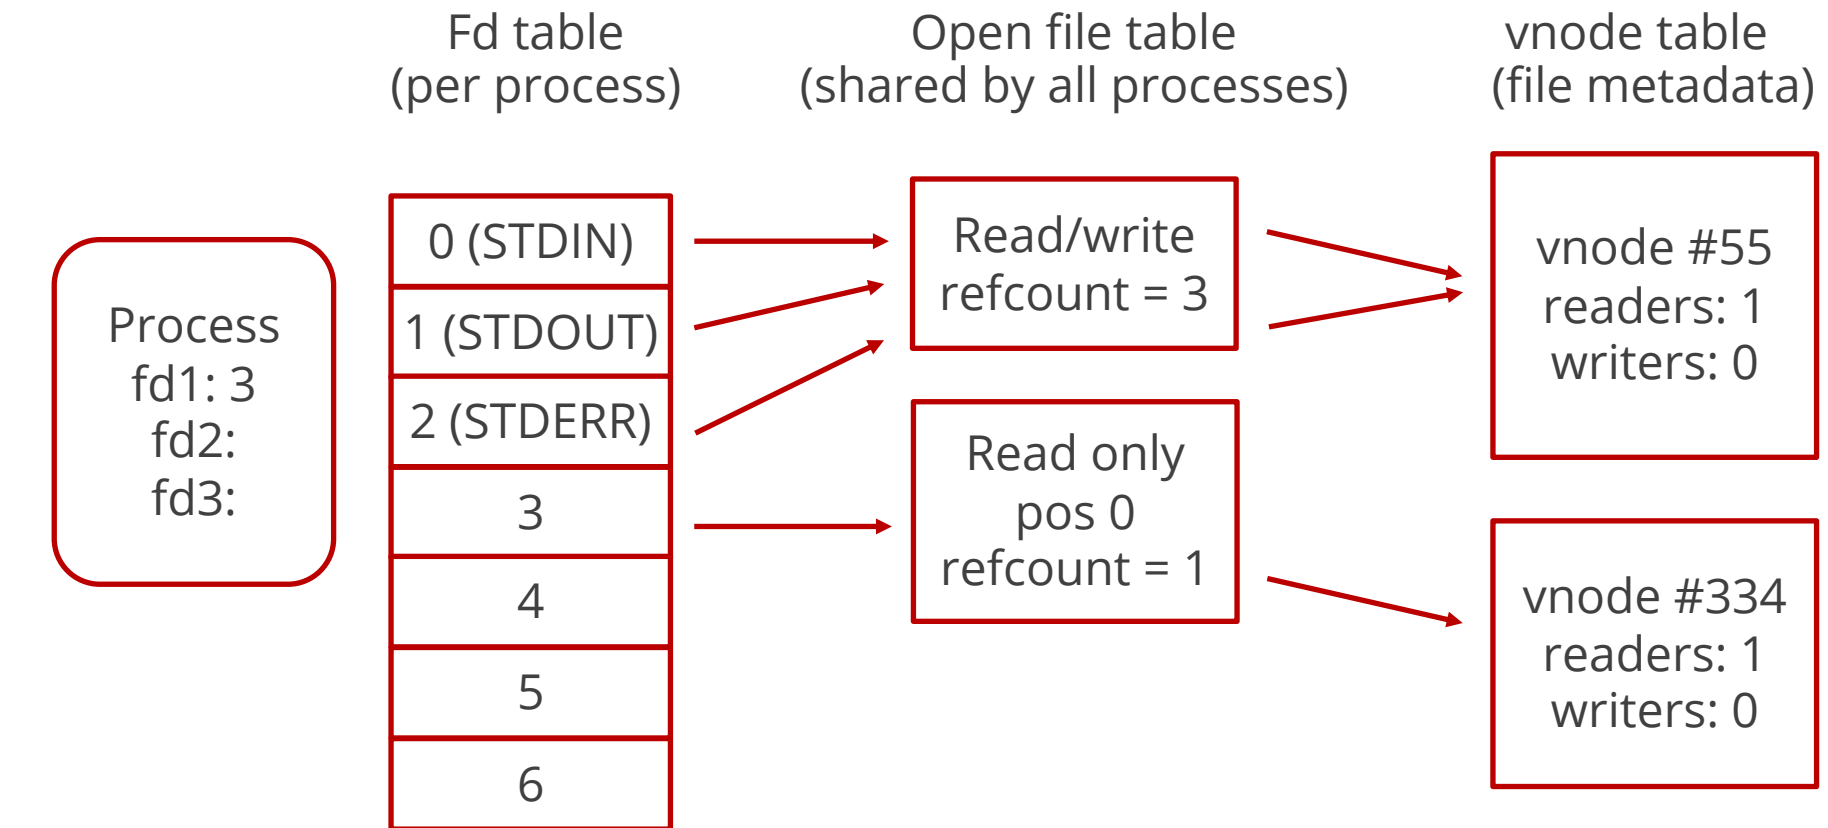
## Open files - example

---

```
char buf[10];  
  
int fd1 = open("foo.c", O_RDONLY, 0);  
int fd2 = dup(fd1);  
int fd3 = open("foo.c", O_RDONLY, 0);
```

What's wrong with this example code?

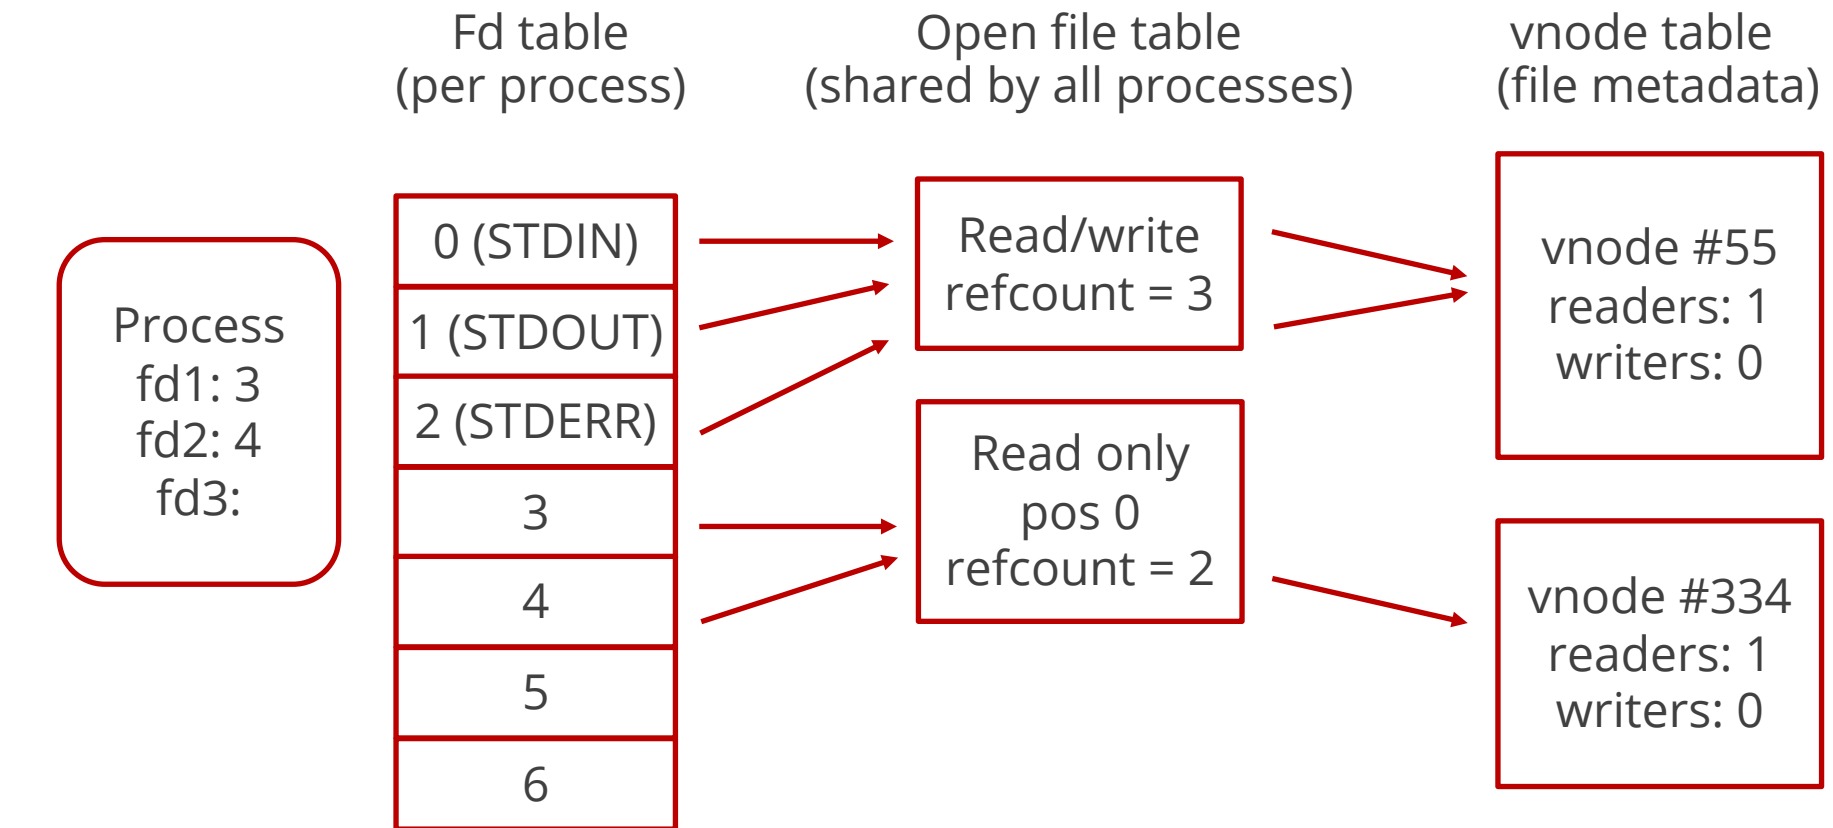
# "Open File" vs "In-Core File"



```
int fd1 = open("foo.c", O_RDONLY, 0);
```

19

# “Open File” vs “In-Core File”

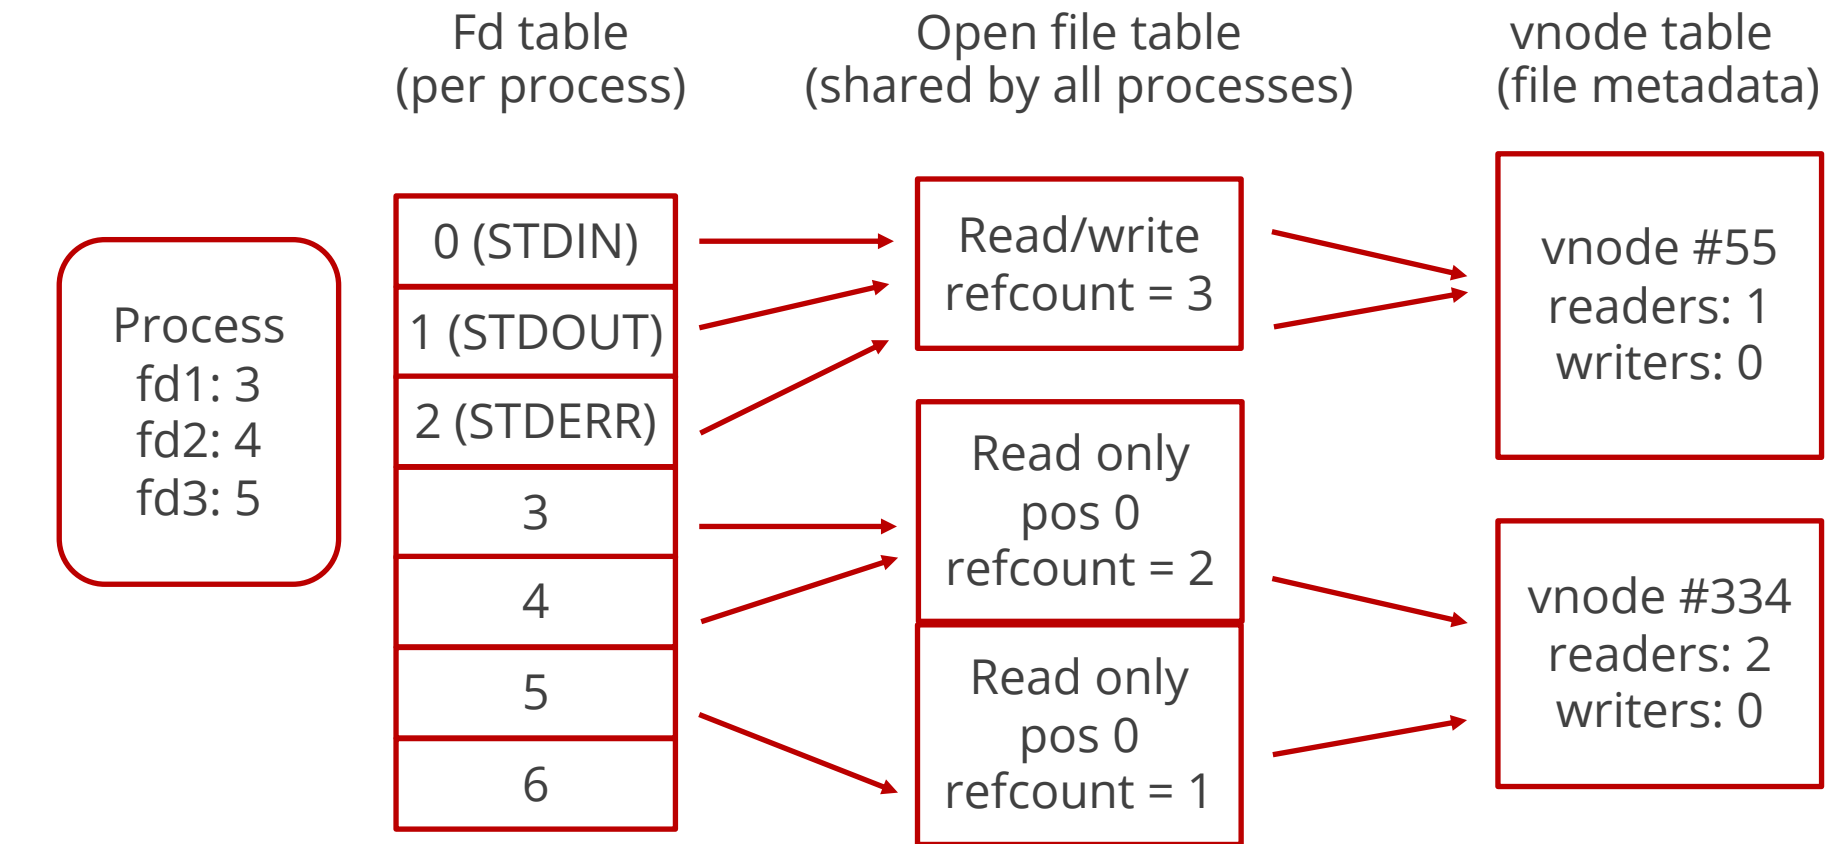


```
int fd2 = dup(fd1);
```

20



# "Open File" vs "In-Core File"



```
21 int fd3 = open("foo.c", O_RDONLY, 0);
```

# Open files - example

---

```
char buf[10];
```

```
int fd1 = open("foo.c", O_RDONLY, 0);
```

```
int fd2 = dup(fd1);
```

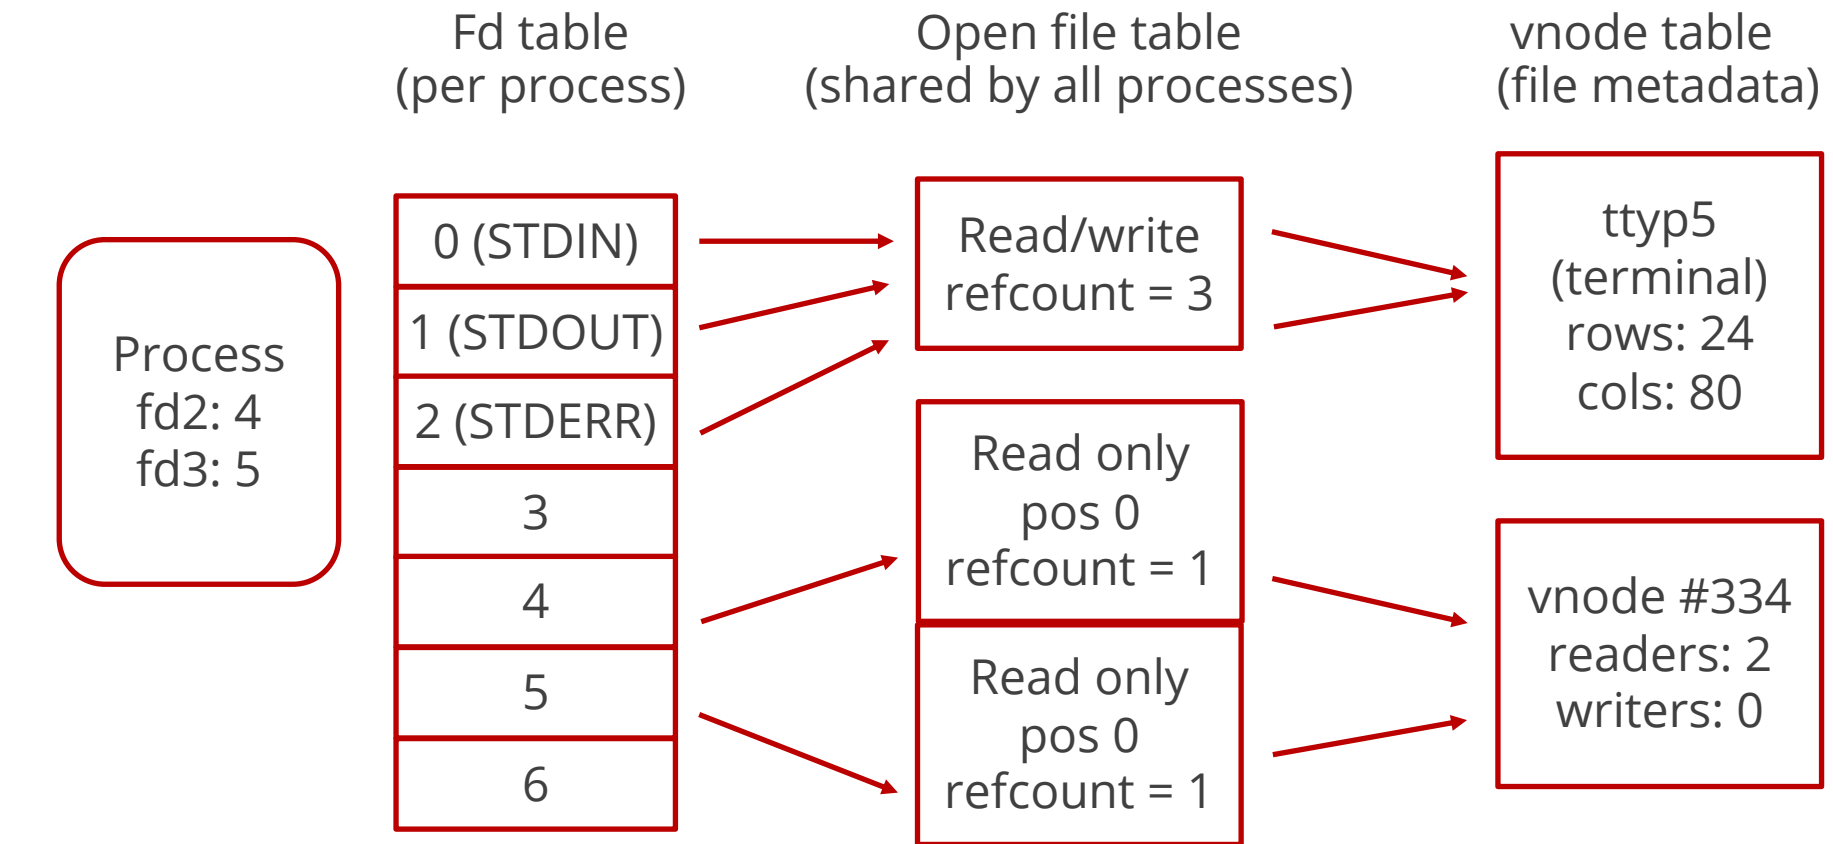
```
int fd3 = open("foo.c", O_RDONLY, 0);
```

```
read(fd1, &buf, sizeof(buf));
```

```
off_t pos2 = lseek(fd2, 0, SEEK_CUR); /* ? */
```

```
off_t pos3 = lseek(fd3, 0, SEEK_CUR); /* ? */
```

# "Open File" vs "In-Core File"



`close(fd1);`  
24

# Let's use file descriptors

---

printf writes to STDOUT. Add to this C code so the second printf writes to "error.txt" using `open` and `close`

```
printf("STDOUT is the command line\n");
```

```
printf("error.txt is STDOUT\n");
```

Challenge: Are there any correctness or performance implications of how `open` works when multithreading?





# What we're covering today: Filesystem interface

---

What are files?

`open` and `fds`

Directories and ownership



# Directory Types

---

## Single-level

- Flat global namespace - only *one* test.c
- Ok for floppy disks (maybe)



# Directory Types

---

## Single-level

- Flat global namespace - only *one* test.c
- Ok for floppy disks (maybe)

## Two-level


- Every user has a directory
- One test.c *per user*
- Typical of early timesharing

# Tree Directories

---

## Directories are special files

- Created with special system calls - `mkdir`
- Format understood + maintained by OS




# Tree Directories

## Directories are special files

- Created with special system calls - `mkdir`
- Format understood + maintained by OS

## Absolute Pathname

- Sequence of directory names
- Starting from "root"
- Ending with a file name
- `/usr/theo`






# Tree Directories

## Current directory (".", cwd)


- "Where I am now" (e.g. /usr/)
- Start of *relative* pathname
  - ./xuesi -> /usr/xuesi
  - ../bin -> /bin/
- Each process has a cwd



# Tree Directories

## Current directory (".", cwd)

- "Where I am now" (e.g. /usr/)
- Start of *relative* pathname
  - ./xuesi -> /usr/xuesi
  - ../bin -> /bin/
- Each process has a cwd



Almost everything is a file description

```
(12:49 sjmcalli@pembroke2 ~) > ls -l /proc/$$/cwd  
lrwxrwxrwx 1 sjmcalli pdl 0 Mar_25 12:46 /proc/43576/cwd -> /h/sjmcalli
```

# Protection - typical baseline

---

Files specifies *owner, group*

- Permissions for owner, permissions for group members
- Permissions for “other” / “world”

Traditional Unix encoding

- `chmod 777 example.txt`
- |   |   |   |   |   |   |   |   |   |   |      |         |
|---|---|---|---|---|---|---|---|---|---|------|---------|
| r | w | x | r | - | x | - | - | x | = | 0751 | (octal) |
| 1 | 1 | 1 | 1 | 0 | 1 | 0 | 0 | 1 |   |      |         |
- V7 Unix: 3 16-bit words specified all permission info
  - permission bits, user #, group #



# What we've covered today: Filesystem interface

---

After the lecture today, you should be able to:

- Explain what a file is and operations that exist on files
- Discuss why we need an open syscall and file descriptors
- Compare different directory types
- Identify owner, group, and permission bits on UNIX files

Next time: How do filesystems enable the file abstraction?