




Filesystems — Interface

MARCH 27, 2025


Sara McAllister
sjmcalli@cs.cmu.edu
saramcallister.github.io

Slides:
saramcallister.github.io/files/2025-03-27-files.pdf



Situating this lecture – Operating Systems (ECE 344)

Date and recording	Lecture Note
1/9	Introduction
1/10	Architecture Support
1/16	Architecture Support
1/17	Architecture Support
1/23	Processes
1/24	Processes
1/30	Threads
1/31	Threads
2/6	Threads, Synchronization (I)
2/7	Synchronization (I), Synchronization (II)
2/13	Synchronization (II)
2/14	Synchronization (II)
2/28	Memory Management (I)
3/6	Memory Management (I)
3/7	Memory Management (I)
3/14	Memory Management (II): Paging
3/20	Memory Management (II): Paging
3/21	Memory Management (III): Replacement



What we've covered so far:

Processes

Threads


Storage Devices



Hard disk drives

Cycle-head-sector (CHS) addressing

Logical block addresses (LBAs)

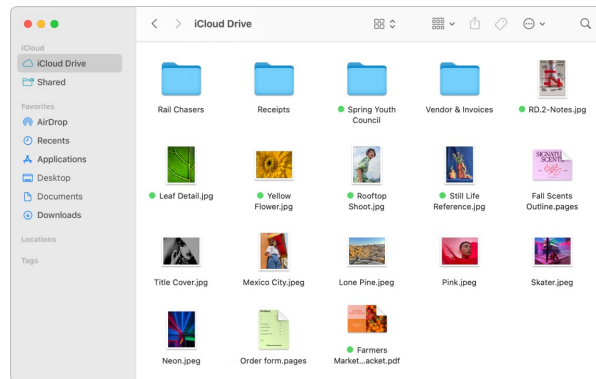


0 1 2 3 4



How can we abstract storage devices?

How do we abstract persistent storage?

Filesystems: a system for organizing, managing, and accessing *files* and *directories* stored on persistent media



Local filesystems (UNIX)



File sharing platforms



Why do we need filesystems?

- Durability across restarts and crashes
- Naming and organization
- Sharing data between processes and users



What makes filesystems challenging?

- Crash recovery (next lecture)
- **Performance + concurrency**
- **Sharing + security**



What we're covering today: Filesystem interface

What are files?

`open` and `fds`

Directories and ownership



What are files?

A file is a sequence of bytes, logically grouped together

- Allow users/ processes to access all files the same way

In UNIX, “Everything is a file”

- `/dev/sda2` (disk partition)
- `/dev/tty2` (terminal)
- `/proc/$$/status` (current process status)

Typical file metadata

Name - 14 characters? 8.3? 255?

Identifier - “file number” (usually internal)

Size - two meanings (next lecture)

Protection - Who can access?

Time, date, last modification

```
struct stat {  
    dev_t      st_dev;          /* ID of device containing file */  
    ino_t      st_ino;          /* Inode number */  
    mode_t     st_mode;         /* File type and mode */  
    nlink_t    st_nlink;        /* Number of hard links */  
    uid_t      st_uid;          /* User ID of owner */  
    gid_t      st_gid;          /* Group ID of owner */  
    dev_t      st_rdev;         /* Device ID (if special file) */  
    off_t      st_size;         /* Total size, in bytes */  
    blksize_t  st_blksize;      /* Block size for filesystem I/O */  
    blkcnt_t   st_blocks;       /* Number of 512B blocks allocated */  
    struct timespec st_atim;    /* Time of last access */  
    struct timespec st_mtim;    /* Time of last modification */  
    struct timespec st_ctim;    /* Time of last status change */  
};
```



Operations on Files

Write, Read - often via position pointer/ "cursor"

Seek - adjust position pointer for next access

Append - write at end of file (implicit synchronization)

Rename - change name of file inside directory and
(maybe) move a file between 2 directories



I/O to a file

Need to read and write to a file

So how do you read and write to one?

Try 1: `read("example.txt", buffer, num_bytes);`

- What if the file is large?

I/O to a file

Need to read and write to a file

So how do you read and write to one?

Try 1: `read("example.txt", buffer, num_bytes);`

- What if the file is large?

Try 2: `read("example.txt", buffer, num_bytes, start_loc);`

- What is notably inefficient about this?
- What's the solution?



What we're covering today: Filesystem interface

What are files?

`open` and `fds`

Directories and ownership



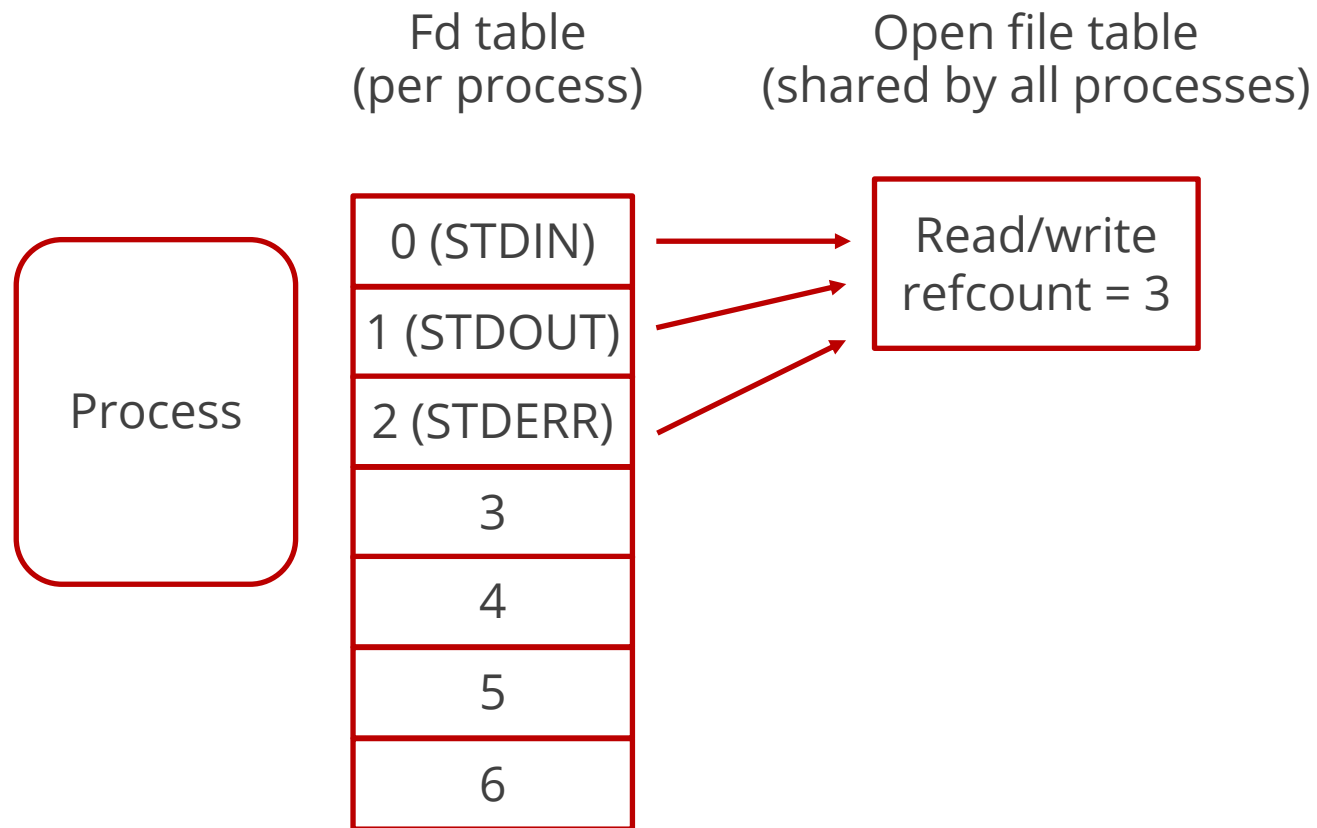
Solution: Open-file state

Add an `open` operation + “state”

“Open-file” structure (file descriptor) stores

- Filesystem / partition
- Filesystem-relative file number
- Operations allowed: eg read vs write
- Cursor position

Open files (UNIX model)





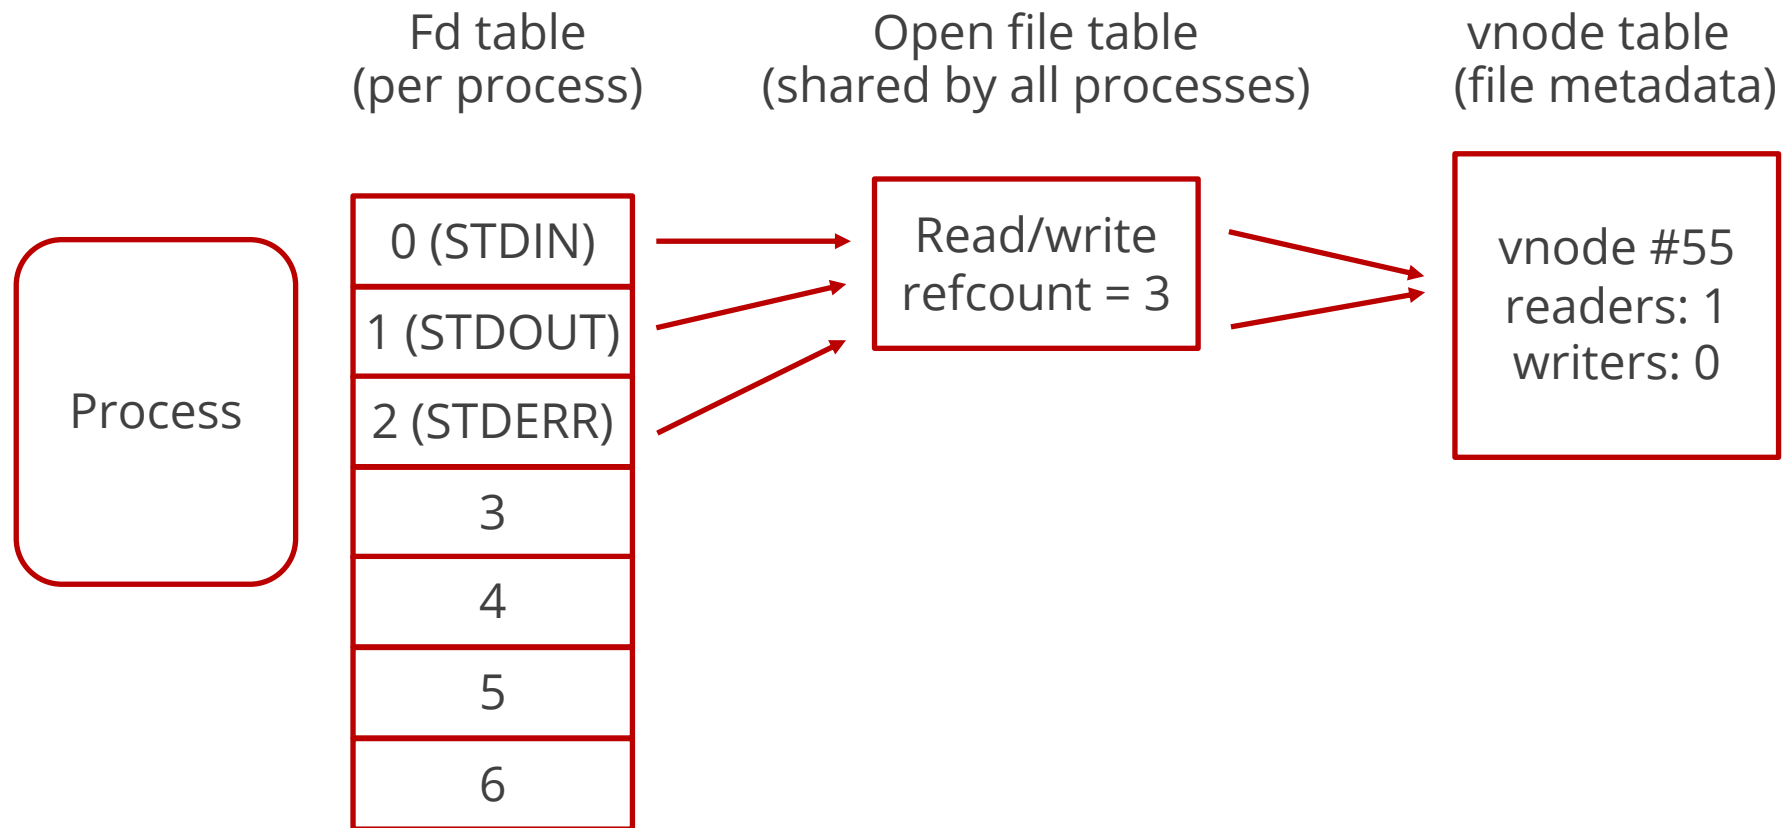
Open files (Unix model)

“In-core” file state - avoid going to disk repeatedly

- Mirror of on-disk structure (File number, size, permissions, modification time, ...)
- Any other filesystem specific information

Shared when file is opened multiple times

Open files (UNIX model)

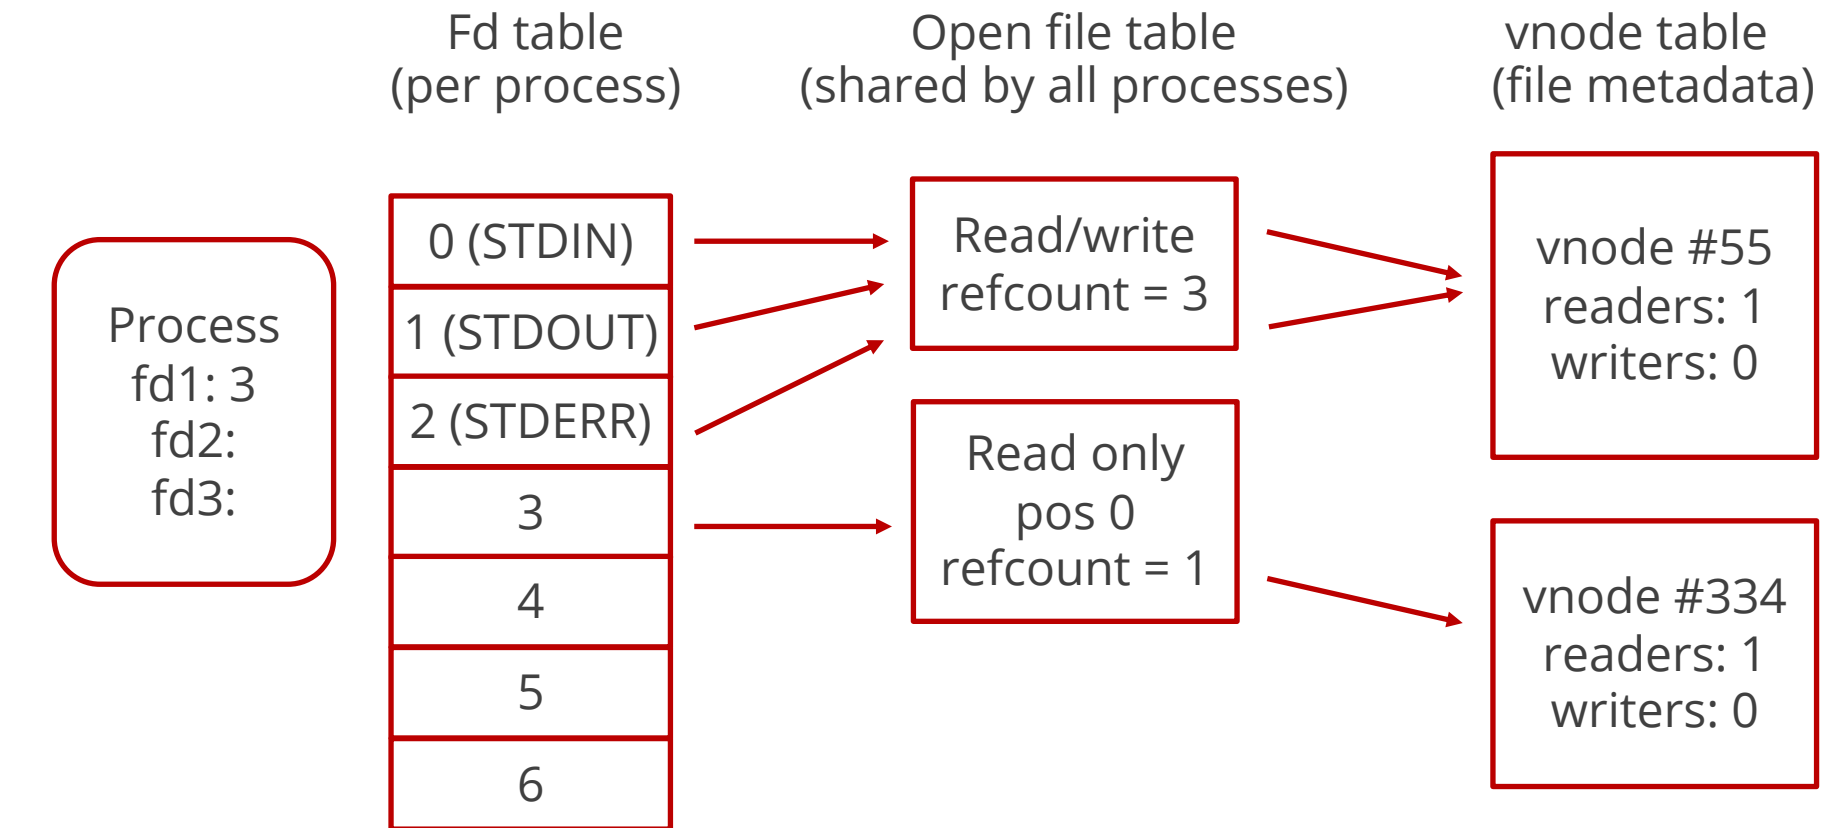


Open files - example

```
char buf[10];  
  
int fd1 = open("foo.c", O_RDONLY, 0);  
int fd2 = dup(fd1);  
int fd3 = open("foo.c", O_RDONLY, 0);
```

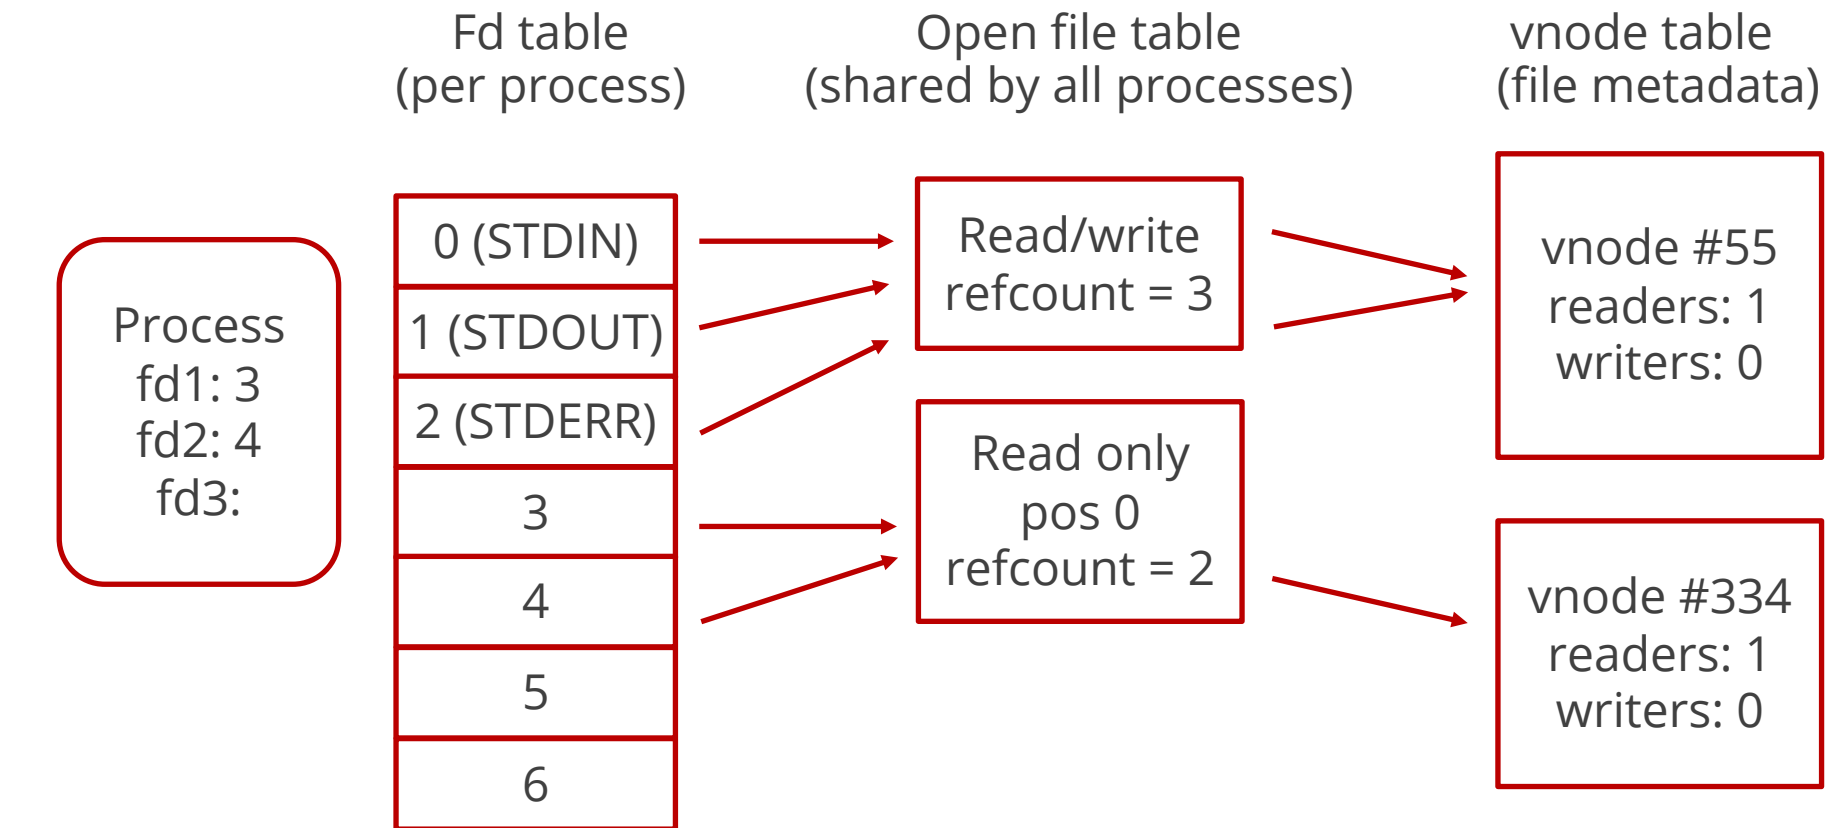
What's wrong with this example code?

"Open File" vs "In-Core File"



```
int fd1 = open("foo.c", O_RDONLY, 0);
```

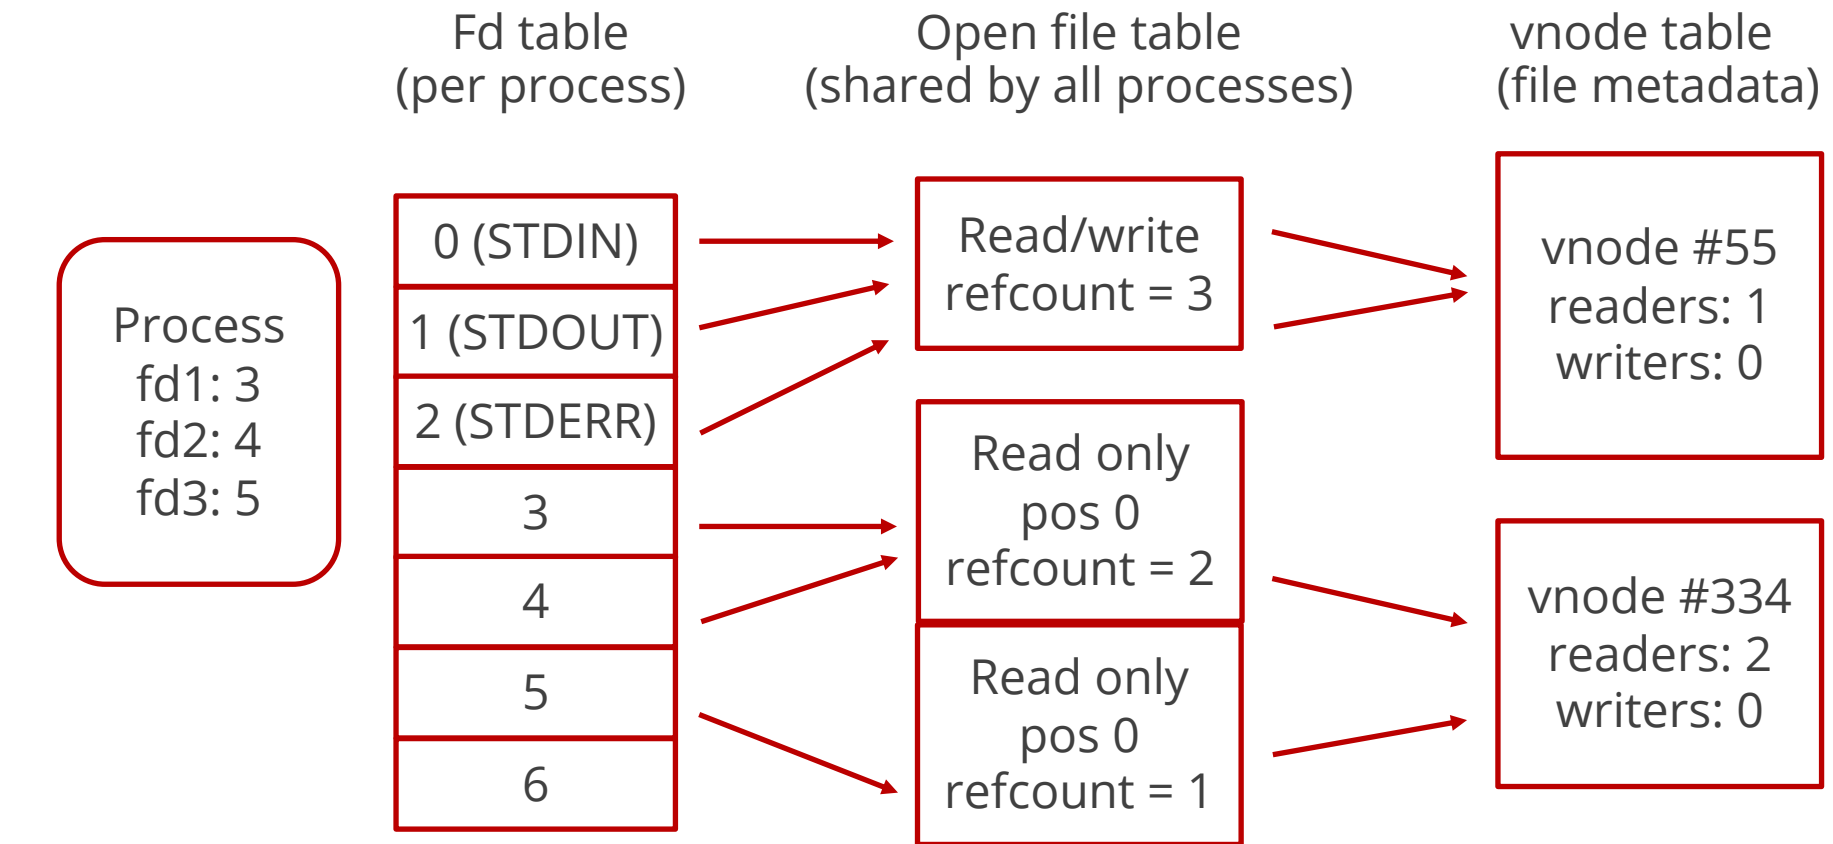
"Open File" vs "In-Core File"



```
int fd2 = dup(fd1);
```

20

"Open File" vs "In-Core File"



21 `int fd3 = open("foo.c", O_RDONLY, 0);`

Open files - example

```
char buf[10];
```

```
int fd1 = open("foo.c", O_RDONLY, 0);
```

```
int fd2 = dup(fd1);
```

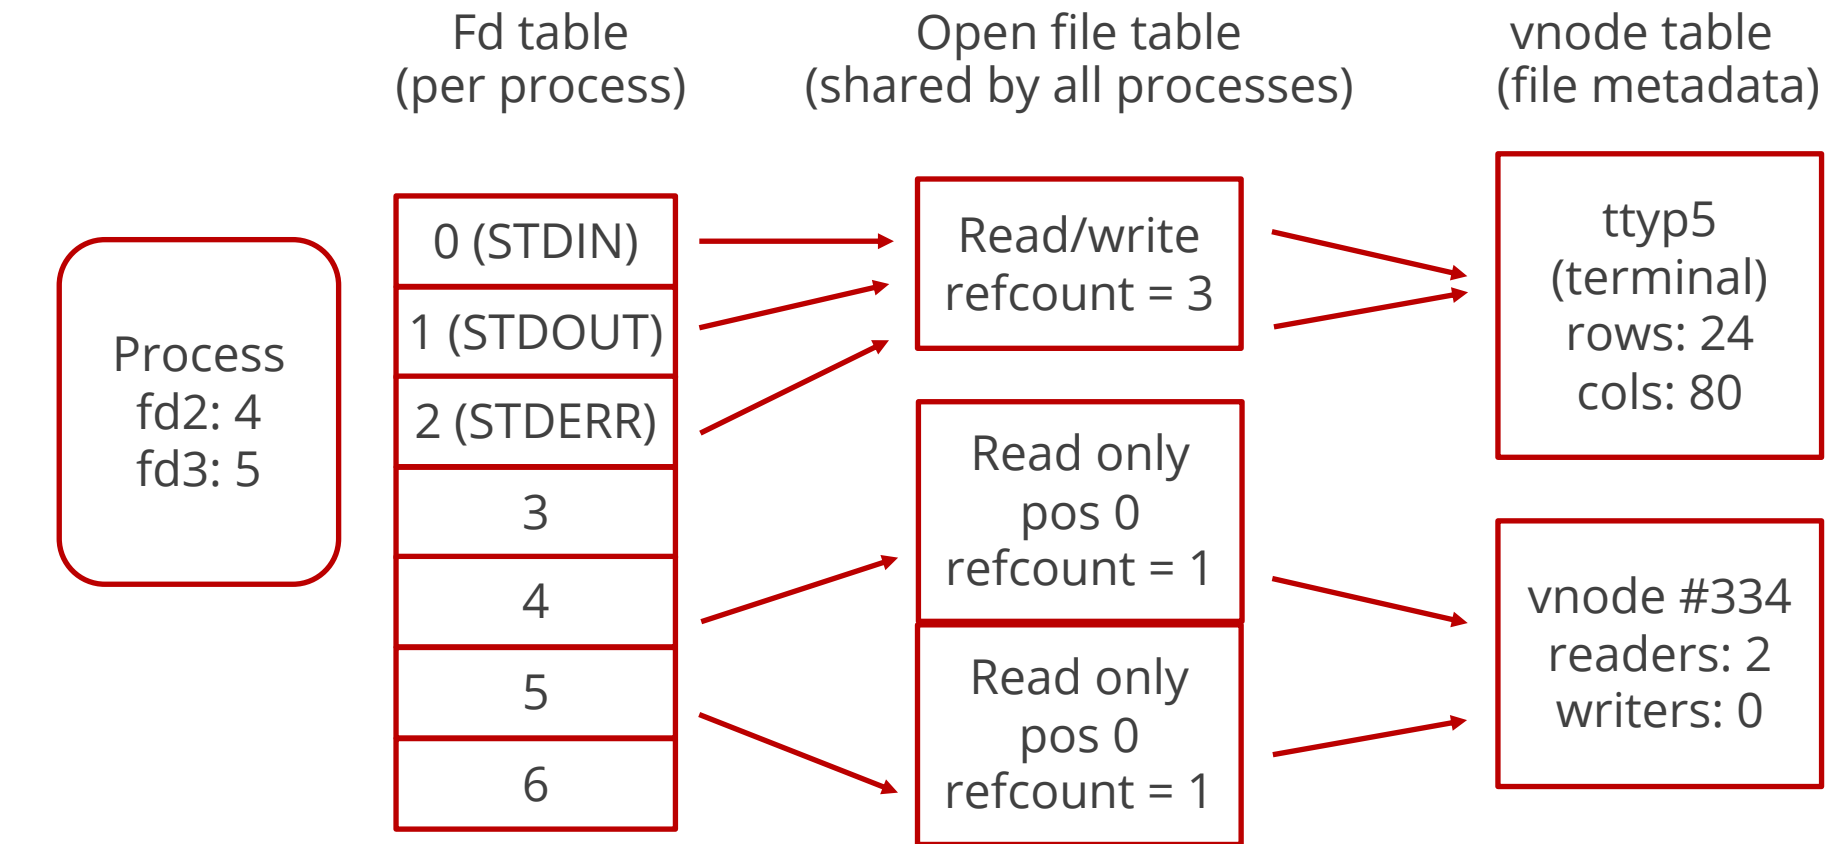
```
int fd3 = open("foo.c", O_RDONLY, 0);
```

```
read(fd1, &buf, sizeof(buf));
```

```
off_t pos2 = lseek(fd2, 0, SEEK_CUR); /* ? */
```

```
off_t pos3 = lseek(fd3, 0, SEEK_CUR); /* ? */
```

"Open File" vs "In-Core File"



`close(fd1);`
24

Let's use file descriptors

printf writes to STDOUT. Add to this C code so the second printf writes to "error.txt" using `open` and `close`

```
printf("STDOUT is the command line\n");
```

```
printf("error.txt is STDOUT\n");
```

Challenge: Are there any correctness or performance implications of how `open` works when multithreading?



What we're covering today: Filesystem interface

What are files?

`open` and `fds`

Directories and ownership



Directory Types

Single-level

- Flat global namespace - only *one* test.c
- Ok for floppy disks (maybe)



Directory Types

Single-level

- Flat global namespace - only *one* test.c
- Ok for floppy disks (maybe)


Two-level

- Every user has a directory
- One test.c *per user*
- Typical of early timesharing

Tree Directories

Directories are special files

- Created with special system calls - `mkdir`
- Format understood + maintained by OS




Tree Directories

Directories are special files

- Created with special system calls - `mkdir`
- Format understood + maintained by OS

Absolute Pathname


- Sequence of directory names
- Starting from "root"
- Ending with a file name
- `/usr/theo`



Tree Directories

Current directory (".", cwd)


- "Where I am now" (e.g. /usr/)
- Start of *relative* pathname
 - ./xuesi -> /usr/xuesi
 - ../bin -> /bin/
- Each process has a cwd



Tree Directories

Current directory (".", cwd)

- "Where I am now" (e.g. /usr/)
- Start of *relative* pathname
 - ./xuesi -> /usr/xuesi
 - ../bin -> /bin/
- Each process has a cwd



Almost everything is a file description

```
(12:49 sjmcalli@pembroke2 ~) > ls -l /proc/$$/cwd  
lrwxrwxrwx 1 sjmcalli pdl 0 Mar_25 12:46 /proc/43576/cwd -> /h/sjmcalli
```

Protection - typical baseline

Files specifies *owner, group*

- Permissions for owner, permissions for group members
- Permissions for “other” / “world”

Traditional Unix encoding

- `chmod 777 example.txt`
- | | | | | | | | | | | | | |
|---|---|---|--|---|---|---|--|---|---|---|---|--------------|
| r | w | x | | r | - | x | | - | - | x | = | 0751 (octal) |
| 1 | 1 | 1 | | 1 | 0 | 1 | | 0 | 0 | 1 | | |
- V7 Unix: 3 16-bit words specified all permission info
 - permission bits, user #, group #



What we've covered today: Filesystem interface

After the lecture today, you should be able to:

- Explain what a file is and operations that exist on files
- Discuss why we need an open syscall and file descriptors
- Compare different directory types

Next time: How do filesystems enable the file abstraction?