

Sara McAllister

PhD Candidate

Carnegie Mellon University

✉ sjmcalli@cs.cmu.edu | 🏠 saramcallister.github.io | 🎓 saramcallister

Research Interests

I research sustainable data centers from a computer systems perspective, and I am particularly interested in caching and storage systems. My work includes a focus on improving efficiency and sustainability through hardware-software co-design and grounding design choices in mathematical modeling. My work has appeared at OSDI and SOSP, including receiving a Best Paper Award at SOSP 2021. I am a 2021 NDSEG fellow and a 2023 EECS Rising Star. I also strive to increase inclusion in computer science, including by creating a DEI course for CS PhD students. Due to these efforts, I was awarded CMU's Graduate Student Service Award in 2022 and a Best Paper Award at SIGCSE 2023.

Education

Carnegie Mellon University

PHD IN COMPUTER SCIENCE, ADVISORS: NATHAN BECKMANN AND GREG GANGER

Pittsburgh, PA

Aug 2019. - Summer 2025 (Expected)

Carnegie Mellon University

MASTERS IN COMPUTER SCIENCE RESEARCH

Pittsburgh, PA

Aug 2019. - May 2022

Harvey Mudd College

B.S. IN COMPUTER SCIENCE, GRADUATED WITH HIGH DISTINCTION

Claremont, CA

Aug. 2015 - May 2019

Honors and Awards

2025 Siebel Scholar, for outstanding academic performance and leadership

2024

2023 Rising Star, in Electrical Engineering & Computer Science (EECS)

2023

NDSEG Exemplary Poster Presentation, in computer and computational sciences at NDSEG fellows conference

2023

SIGCSE Best Paper Award, for CS-JEDI paper

2023

CMU Graduate Student Service Award, for the development of 15-996 CS-JEDI

2022

SOSP Best Paper Award, for Kangaroo paper

2021

NDSEG Graduate Fellowship, DoD sponsored 3-year fellowship

2021

NSF Graduate Research Fellowship (GRFP), NSF sponsored 3-year fellowship

2021

Harvey Mudd Class of '94 Award, for an outstanding CS graduate in coursework, research, and service

2019

Harvey Mudd Computer Science Departmental Honors

2019

Harvey Mudd Clinic Team Award, for outstanding performance on an industry-sponsored team capstone project

2019

CRA Outstanding Undergraduate Researcher Award, Honorable Mention

2019

Publications

FairyWREN: A Sustainable Cache for Emerging Write-Read-Erase Flash Interfaces

OSDI 2024

[Sara McAllister](#), Yucong Wang, Benjamin Berg, Daniel S. Berger, George Amvrosiadis, Nathan Beckmann, Gregory R. Ganger

Acceptance Rate: 18%

A Call for Research on Storage Emissions

HotCarbon 2024

[Sara McAllister](#), Fiodar Kazhamiaka, Daniel S. Berger, Rodrigo Fonseca, Kali Frost, Aaron Ogus, Maneesh Sah, Ricardo Bianchini, George Amvrosiadis, Nathan Beckmann, Gregory R. Ganger

Acceptance Rate: 46%

DéjàVu: KV-cache Streaming for Fast, Fault-tolerant Generative LLM Serving

ICML 2024

Fonteini Strati, [Sara McAllister](#), Amar Phanishayee, Jakub Tarnawski, Ana Klimovic

Acceptance Rate: 27.5%

Towards Understanding the Carbon Impact in End-to-end Sensing Pipelines

HotEthics 2024

Harsh Desai*, [Sara McAllister](#)*, Nathan Beckmann, Brandon Lucia (* = co-first author)

CS-JEDI: Required DEI Education, by CS PhD Students, for CS PhD Students

🏆 SIGCSE 2023

Bailey Flanigan, Ananya Joshi, [Sara McAllister](#), Catalina Vajiac

Acceptance Rate: 35%

Kangaroo: Theory and Practice of Caching Billions of Tiny Objects on Flash

ACM TOS

[Sara McAllister](#), Benjamin Berg, Julian Tutuncu-Macias, Juncheng Yang, Sathya Gunasekar, Jimmy Lu, Daniel S. Berger, Nathan Beckmann, Gregory R. Ganger

August 2022

Kangaroo: Caching Billions of Tiny Objects on Flash

Sara McAllister, Benjamin Berg, Julian Tutuncu-Macias, Juncheng Yang, Sathya Gunasekar, Jimmy Lu, Daniel S. Berger, Nathan Beckmann, Gregory R. Ganger

 SOSP 2021

Acceptance Rate: 16%

External-memory Dictionaries in the Affine and PDAM Models

Michael A. Bender, Alex Conway, Martin Farach-Colton, William Jannen, Yizheng Jiao, Rob Johnson, Eric Knorr, Sara McAllister, Nirjhar Mukherjee, Prashant Pandey, Donald E. Porter, Jun Yuan, Yang Zhan

ACM ToPC

September 2021

The CacheLib Caching Engine: Design and Experiences at Scale

Benjamin Berg, Daniel S. Berger, Sara McAllister, Isaac Grosof, Sathya Gunasekar, Jimmy Lu, Michael Uhlar, Jim Carrig, Nathan Beckmann, Mor Harchol-Balter, Gregory R. Ganger

OSDI 2020

Acceptance Rate: 18%

Small Refinements to DAM Can Have Big Consequences for Data-Structure Design

Michael A. Bender, Alexander Conway, Martin Farach-Colton, William Jannen, Yizheng Jiao, Rob Johnson, Eric Knorr, Sara McAllister, Nirjhar Mukherjee, Prashant Pandey, Donald E. Porter, Jun Yuan, Yang Zhan

SPAA 2019

Acceptance Rate: 40%

Talks

A Call for Research on Storage Emissions

Western Digital (Remote) – Hosted by Toshiki Hirano
HotCarbon

5 Sep 2024

9 July 2024

FairyWREN: A Sustainable Cache for Write-Read-Erase Interfaces

OSDI
PDL Retreat – Presented to a large group of industry attendees
PDL Retreat – Presented to a large group of industry attendees

12 July 2024

7 Nov 2023

7 Nov 2022

Towards Understanding the Carbon Impact in End-to-end Sensing Pipelines

HotEthics – Co-presented with Harsh Desai

29 Apr 2024

Overcoming Write Limitations to achieve Sustainable Flash Caching

AMD (Remote) – Research and Advance Development (RAD) and Xilinx Labs
Salesforce (Remote) – Database Reading Group
UC Berkeley – Hosted by Natacha Crooks
Stanford – Hosted by Keith Winstein
UC Santa Cruz – Hosted by Andrew Quinn
McGill (Remote) – Hosted by Oana Balmau
Microsoft Pittsburgh – Hosted by Jeff Butler
MIT – Hosted by Frans Kaashoek
NDSEG 2021 Fellows Conference – Recieved best poster-presentation award
University of Toronto – Hosted by Bianca Schroeder

29 Mar 2024

27 Mar 2024

25 Jan 2024

24 Jan 2024

11 Jan 2024

16 Nov 2023

2 Nov 2023

10 Oct 2023

31 July 2023

20 Mar 2023

Scaling the bandwidth-per-TB wall with Declarative Storage Interfaces

PDL Retreat – Presented to a large group of industry attendees

6 Nov 2023

CS-JEDI: DEI education by PhD students, for PhD students

McGill (Remote) – Hosted by Oana Balmau

31 Oct 2023

Caching on Flash: Kangaroo and Beyond

Meta (Remote) – Core Data Tech Talk

11 Mar 2022

Kangaroo: Caching Billions of Objects on Flash

Microsoft Research (Remote) – Hosted by Daniel Berger
SOSP (Remote)
Cache@Scale (Remote) – Industry Caching Meetup hosted by Meta

22 Nov 2021

27 Oct 2021

4 Mar 2021

Building a Stronger, More Just Academic Community Through Mandatory Anti-bias Learning

University of Pittsburgh Diversity Forum (Remote) – Co-presented w/ Bailey Flanigan and Catalina Vajiac

28 July 2021

Teaching

Carnegie Mellon University

Storage Systems (15-746/18-746)

TA, Fall 2023

Parallel Computer Architecture and Programming (15-418/618)

TA, Spring 2022

Diversity, Equity, and Inclusion in Computer Science and Society (15-996)

Co-Creator and TA, Spring 2021

Harvey Mudd College

Programming Languages (CS131)

Grader and Tutor, Spring 2019

Introduction to Computer Systems (CS105)

Grader and Tutor, Fall 2018

Introduction to Computer Systems (CS105)

Grader and Tutor, Spring 2018

Data Structures and Programming Development (CS70)

Grader and Tutor, Fall 2017

Principles of Computer Science (CS60)

Grader and Tutor, Spring 2017

Introduction to Biology and Computer Science (CS5 Green)

Grader and Tutor, Fall 2016

Guest Lecturer

Storage Systems – Overcoming Flash’s Write Limitations to Achieve Sustainable Caching (CMU 15/18-746)

Fall 2023

Graduate Computer Architecture – Sustainable Computing (CMU 15-740)

Fall 2023

Computer Systems – Kangaroo: Caching Billions of Tiny Objects on Flash (CMU 18-213/613)

Fall 2022

Data Center Computing – Kangaroo Discussion (CMU 18-847C)

Spring 2022

CS-JEDI – Panel on Allyship (CMU 15-996)

Spring 2022

Computer Systems – Kangaroo: Caching Billions of Tiny Objects on Flash (CMU 18-213)

Fall 2021

Storage Systems – Kangaroo: Caching Billions of Tiny Objects on Flash (CMU 18-746)

Fall 2021

Mentoring

Lucy Wang. CMU ECE undergraduate student

Spring 2024 - Present

Suhas Thalanki. CMU computational data science masters capstone

Spring 2024 - Present

Sriya Ravi. CMU computational data science masters capstone

Spring 2024 - Present

Yu Liu. CMU computational data science masters capstone

Spring 2024 - Present

Sophia (Qingyang) Cao. CMU CS undergraduate student

Fall 2023 - Present

Sarvesh Tandon. CMU ECE masters student

Fall 2023 - Present

Sherry (Yucong) Wang. CMU ECE undergraduate student, After degree: Salesforce

Fall 2022 - Spring 2024

Akshath Karanam. CMU ECE masters student, After degree: Salesforce

Fall 2022

Priyal Suneja. University of Washington CS PhD student

Fall 2021 - Summer 2022

Julian Tutuncu-Macias. CMU CS undergraduate student, After degree: Goldman Sachs

Fall 2019 - Spring 2021

Sheng Xu. CMU CS masters student, After degree: Amazon Web Services

Spring 2020

Karina Mejia. Ontario High School

Summer 2016

Service

Reviewer

Transactions on Storage (TOS)

2024

External Review Committee Member

USENIX Annual Technical Conference (ATC)

2024

Faculty Hiring Committee

Carnegie Mellon University, Computer Science Department

2024

Harvey Mudd College, Computer Science Department

2019

PhD Admissions

Carnegie Mellon University, Computer Science Department

2022

Student Organizer

DEI initiatives in CMU’s CS Department – *Informal Survey, CS-JEDI course, advisor-advisee feedback form*

2020-2023

Parallel Data Lab (PDL) Meeting Coordinator

2021

PhD Orientation Committee – *CMU CS Department’s Introductory Course (IC)*

2020

Community Outreach

Professional Experience

Graduate Research Assistant

ADVISORS: NATHAN BECKMANN AND GREG GANGER

Researched reducing IO to create sustainable caching and storage systems at scale

Carnegie Mellon University

Aug. 2019 - Present

Research Intern

MENTOR: AMAR PHANISHAYEE

Researched serving LLMs more efficiently especially under failure

Microsoft Research

Summer 2022

Research Intern

MENTOR: DANIEL BERGER

Researched in-kernel disaggregated CXL memory solutions

Microsoft Research

Summer 2021

Software Engineering Intern

DATABASE TEAM

Designed, implemented, and rolled out a library to manage MySQL database permissions

Yelp

Summer 2019

Clinic (Capstone) Project

SPONSORED BY PURE STORAGE

Technical lead for team of 4 designing failover mechanisms for NFS VMs running on a two-controller system

Harvey Mudd College

Aug. 2018 - May 2019

Undergraduate Research Assistant

ADVISOR: DON PORTER

Researched theoretical and experimental analysis of write-optimized dictionaries

UNC Chapel Hill

May 2018 - Aug. 2018

Software Engineering Intern

DEVELOPER EXPERIENCE TEAM

Developed a Python library to restart and repair development servers

Facebook

Summer 2017

Research Assistant

ADVISOR: ANNA AHN

Researched three-legged walking and led data analysis of wearable devices

Harvey Mudd College

May. 2016 - Jun. 2017